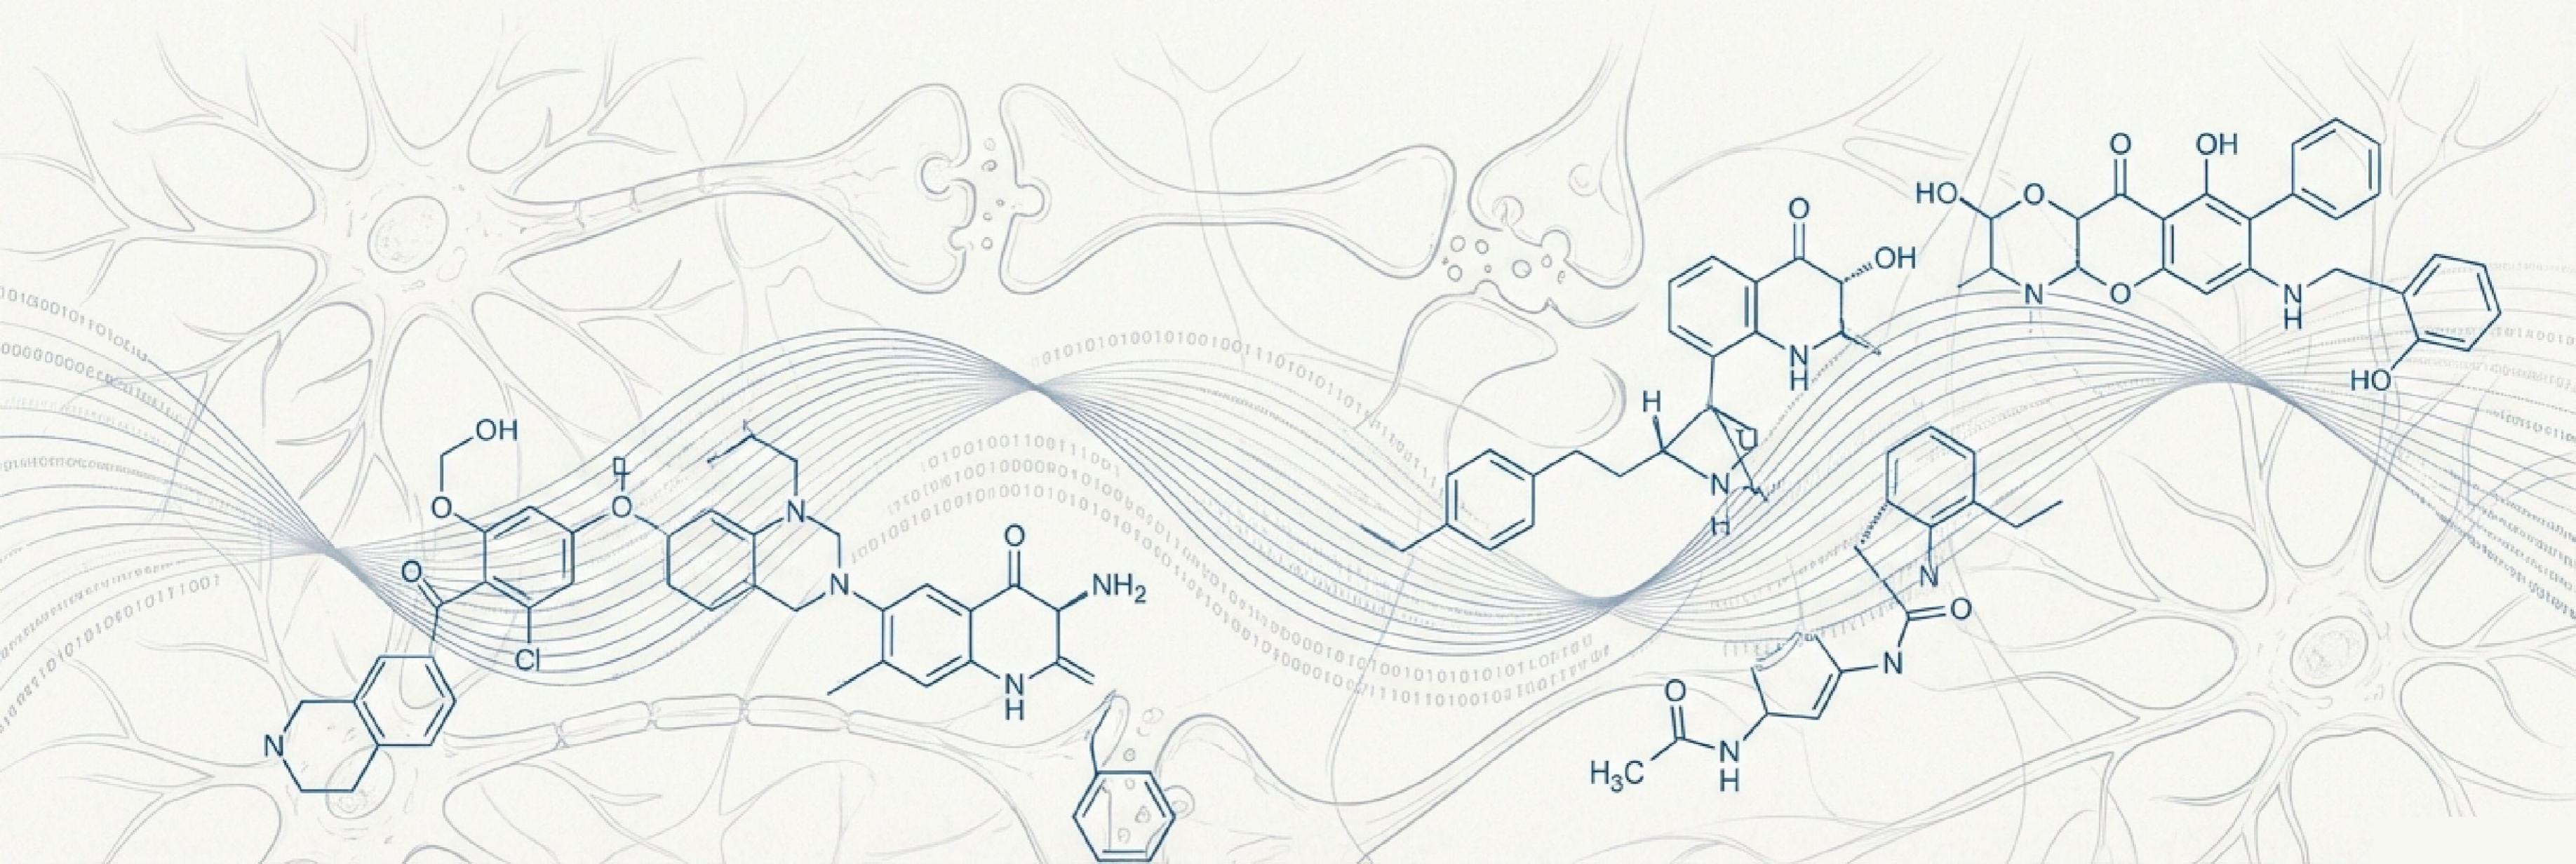


Accelerating Alzheimer's Research with Machine Learning

A QSAR Case Study for Predicting Drug Potency Against Beta-Amyloid



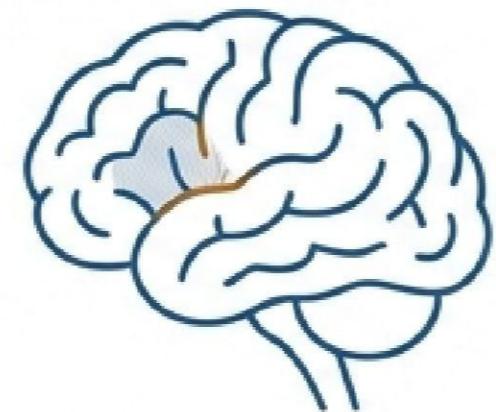
The Immense Challenge of Alzheimer's Disease

Alzheimer's is a progressive neurodegenerative disease that destroys memory and cognitive function. It represents a global health crisis with no current cure.



50 million

people affected globally.



60-70%

of all dementia cases
worldwide.



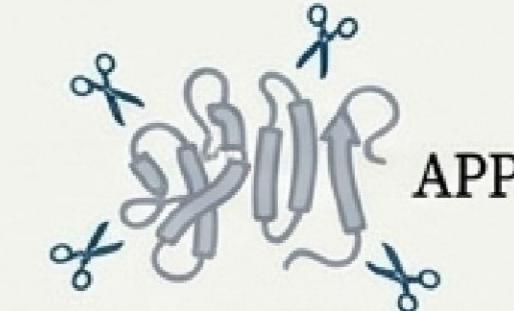
6th

leading cause of death in
the United States.

Pinpointing the Culprit: The Amyloid Hypothesis

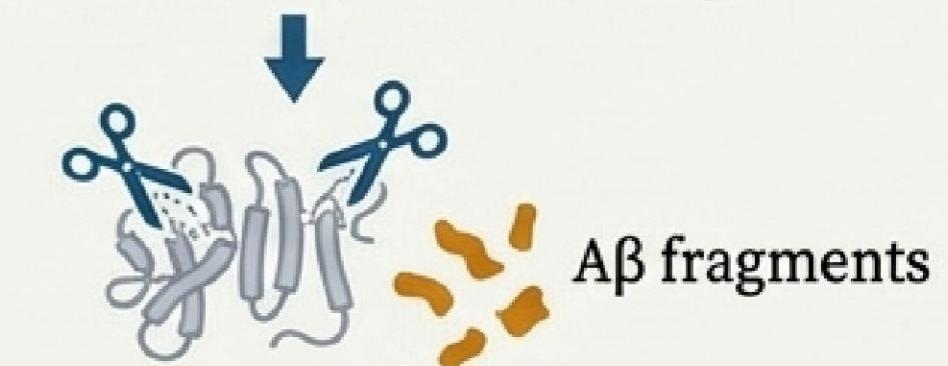
A leading theory points to the Beta-Amyloid A4 protein ($\text{A}\beta$) as a key player. In Alzheimer's, this protein accumulates and forms toxic clumps, leading to a cascade of cellular damage.

1



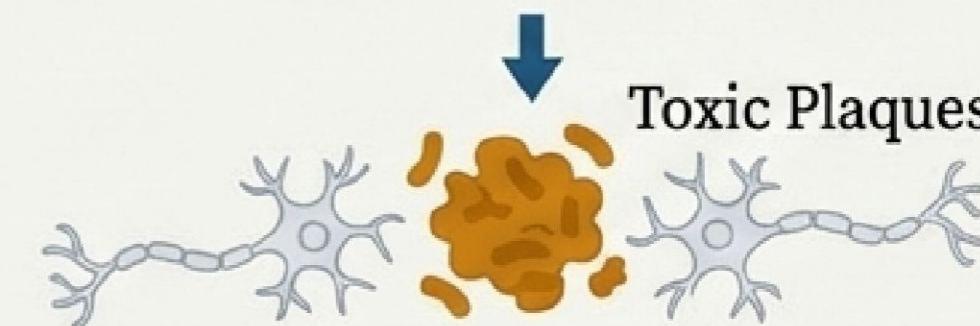
A larger protein (APP) is cut by enzymes.

2



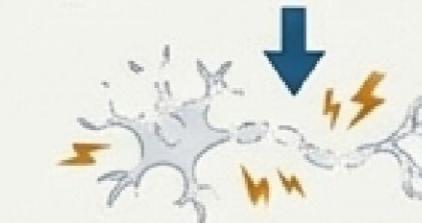
This produces sticky $\text{A}\beta$ fragments.

3



Fragments clump together to form toxic plaques between neurons.

4

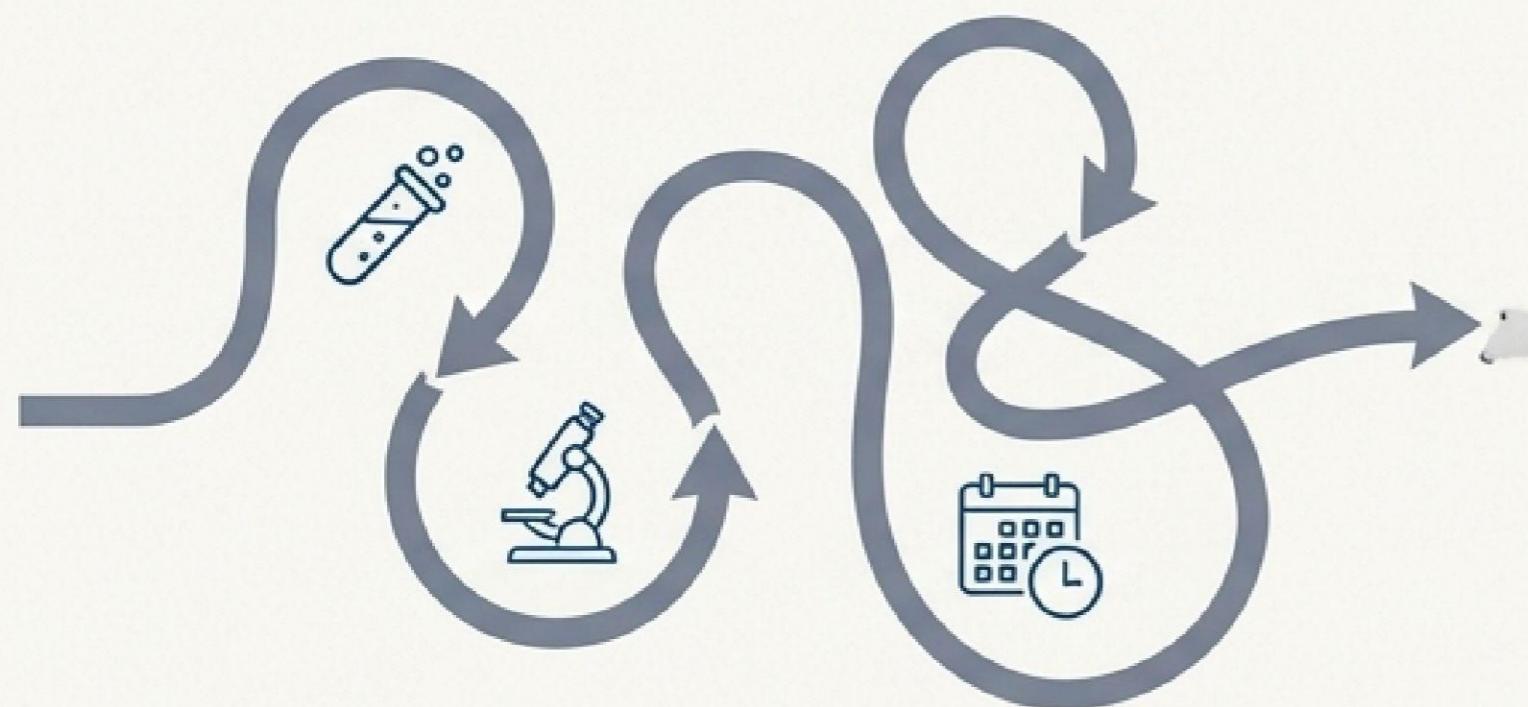


This leads to synaptic dysfunction, chemical compounds (inhibitors) that can stop this process.

Our goal is to find chemical compounds (inhibitors) that can stop this process.

A Faster Path: Predicting Potency with Quantitative Structure-Activity Relationship (QSAR)

The Traditional Approach



Years of slow, iterative lab testing.

Average cost: **\$2.6 billion** per approved drug.

The QSAR Approach



A computational framework built on the hypothesis:
Biological Activity = Function(Chemical Structure)

By learning the relationship between a molecule's structure and its effectiveness, we can use machine learning to predict the potency of millions of untested compounds, saving immense time and resources.

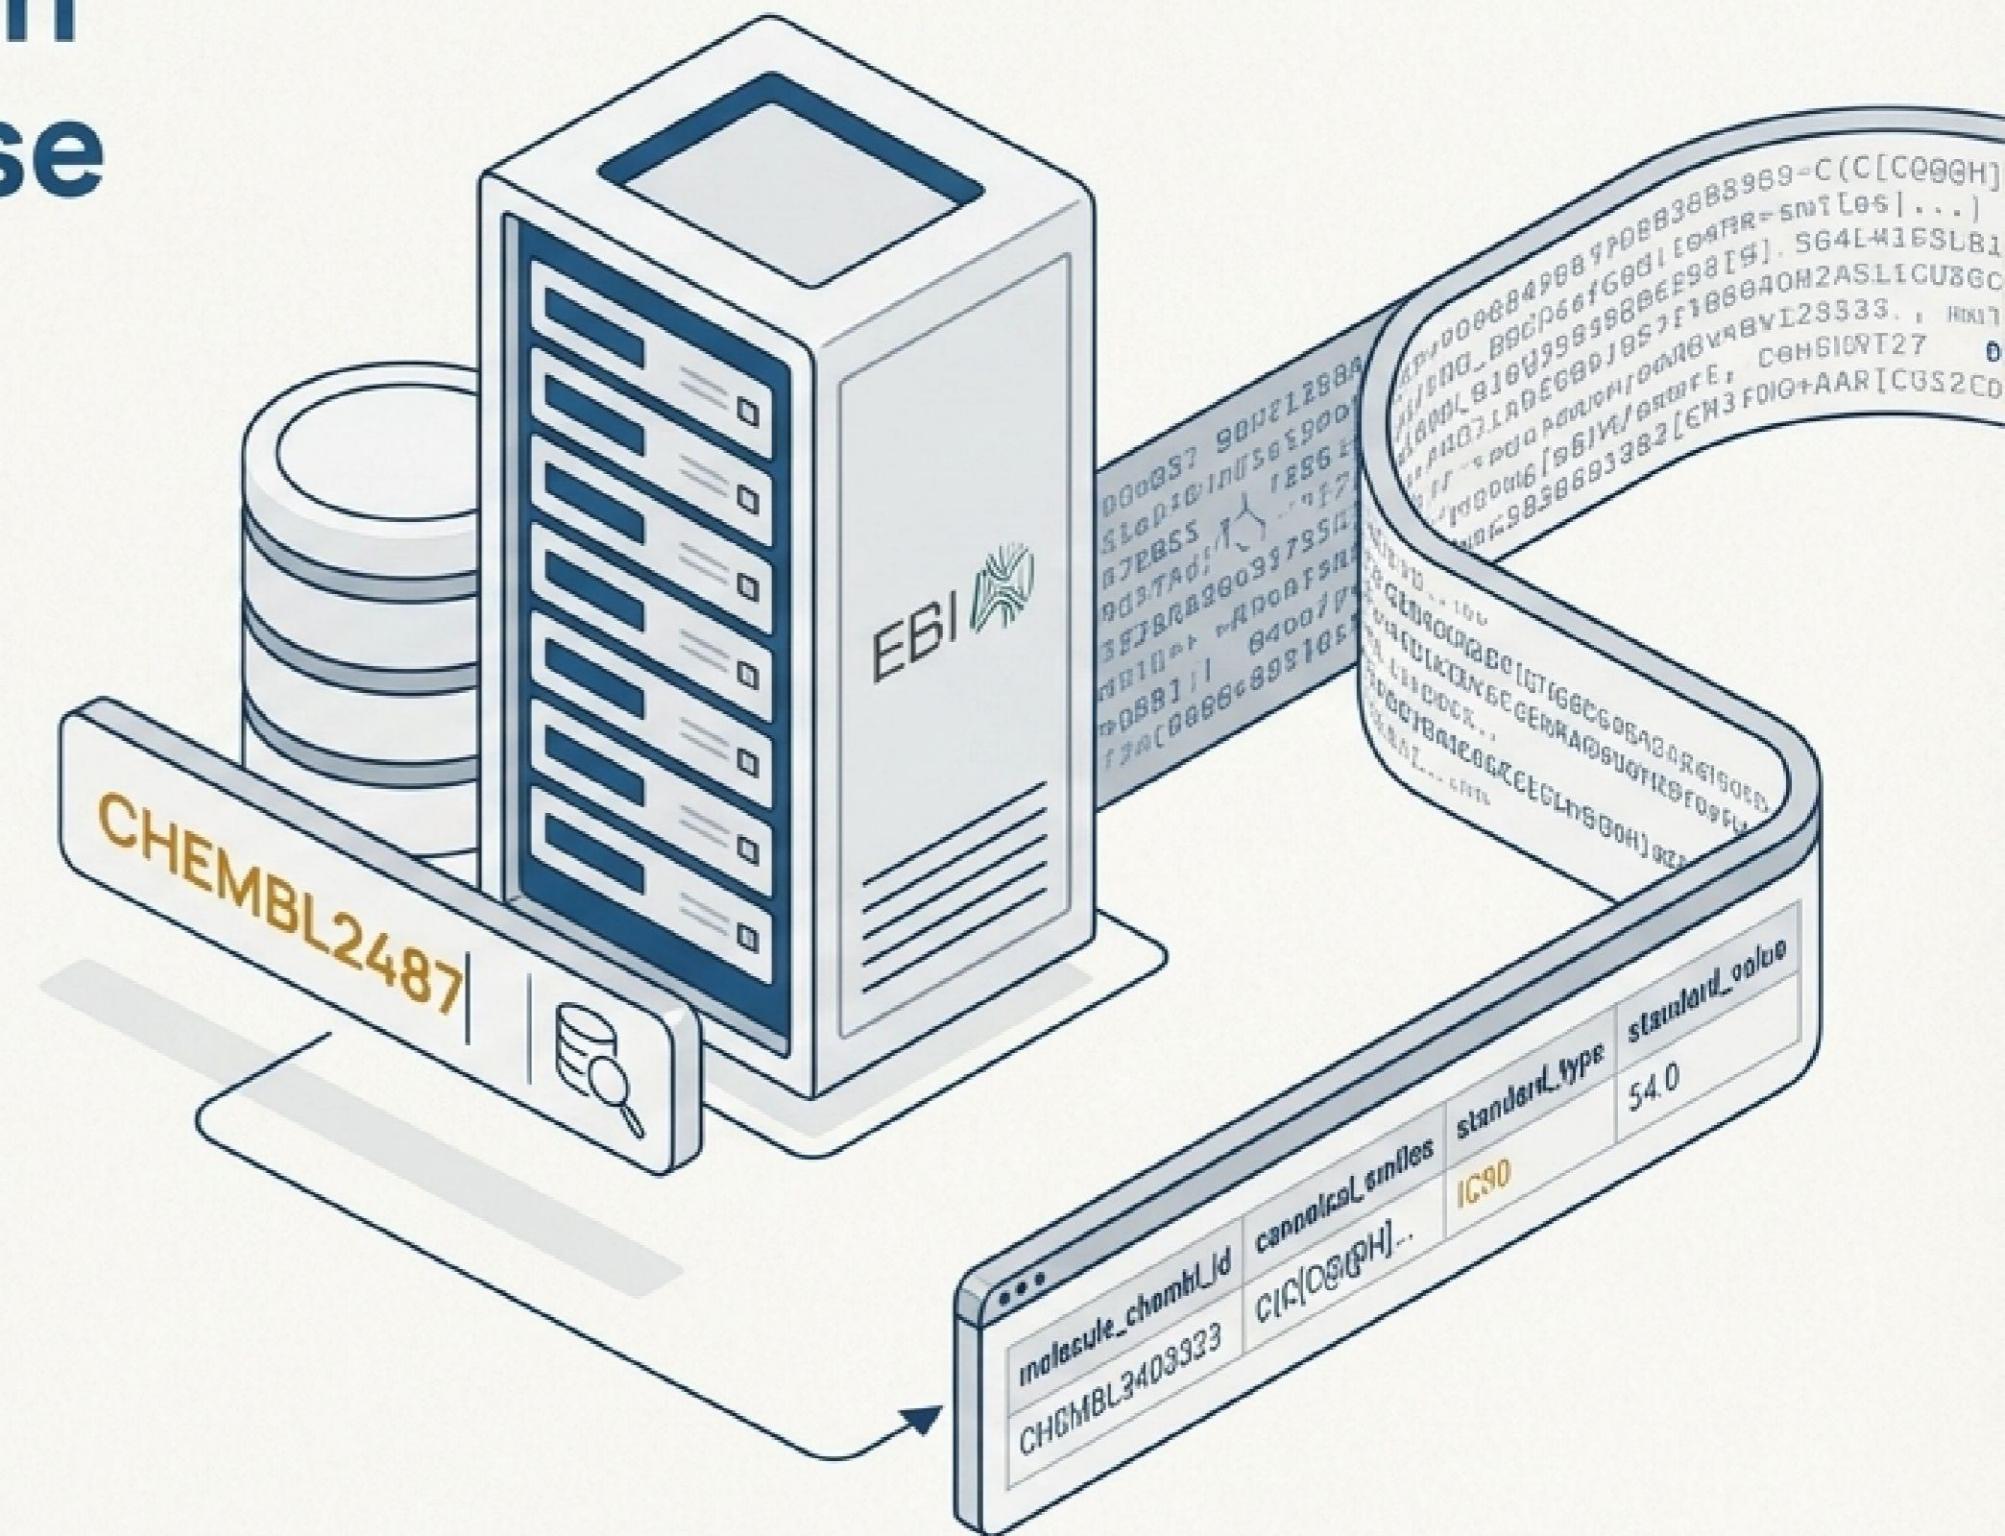
Mining for Ground Truth in the ChEMBL Database

What is ChEMBL?

A massive, open-access database from the European Bioinformatics Institute, containing experimental results for bioactive, drug-like molecules. It is the gold standard for this type of research.

Our Query

- Target Protein: Beta-amyloid A4
- ChEMBL ID: **CHEMBL2487**
- Goal: Retrieve all recorded experiments measuring the potency of chemical inhibitors against this specific target.



Forging a Clean and Consistent Dataset

Raw experimental data is noisy. To train a reliable model, we must standardize and clean it through a multi-step process.

Start with Raw Data

All bioactivity entries for CHEMBL2487.

Filter for IC50

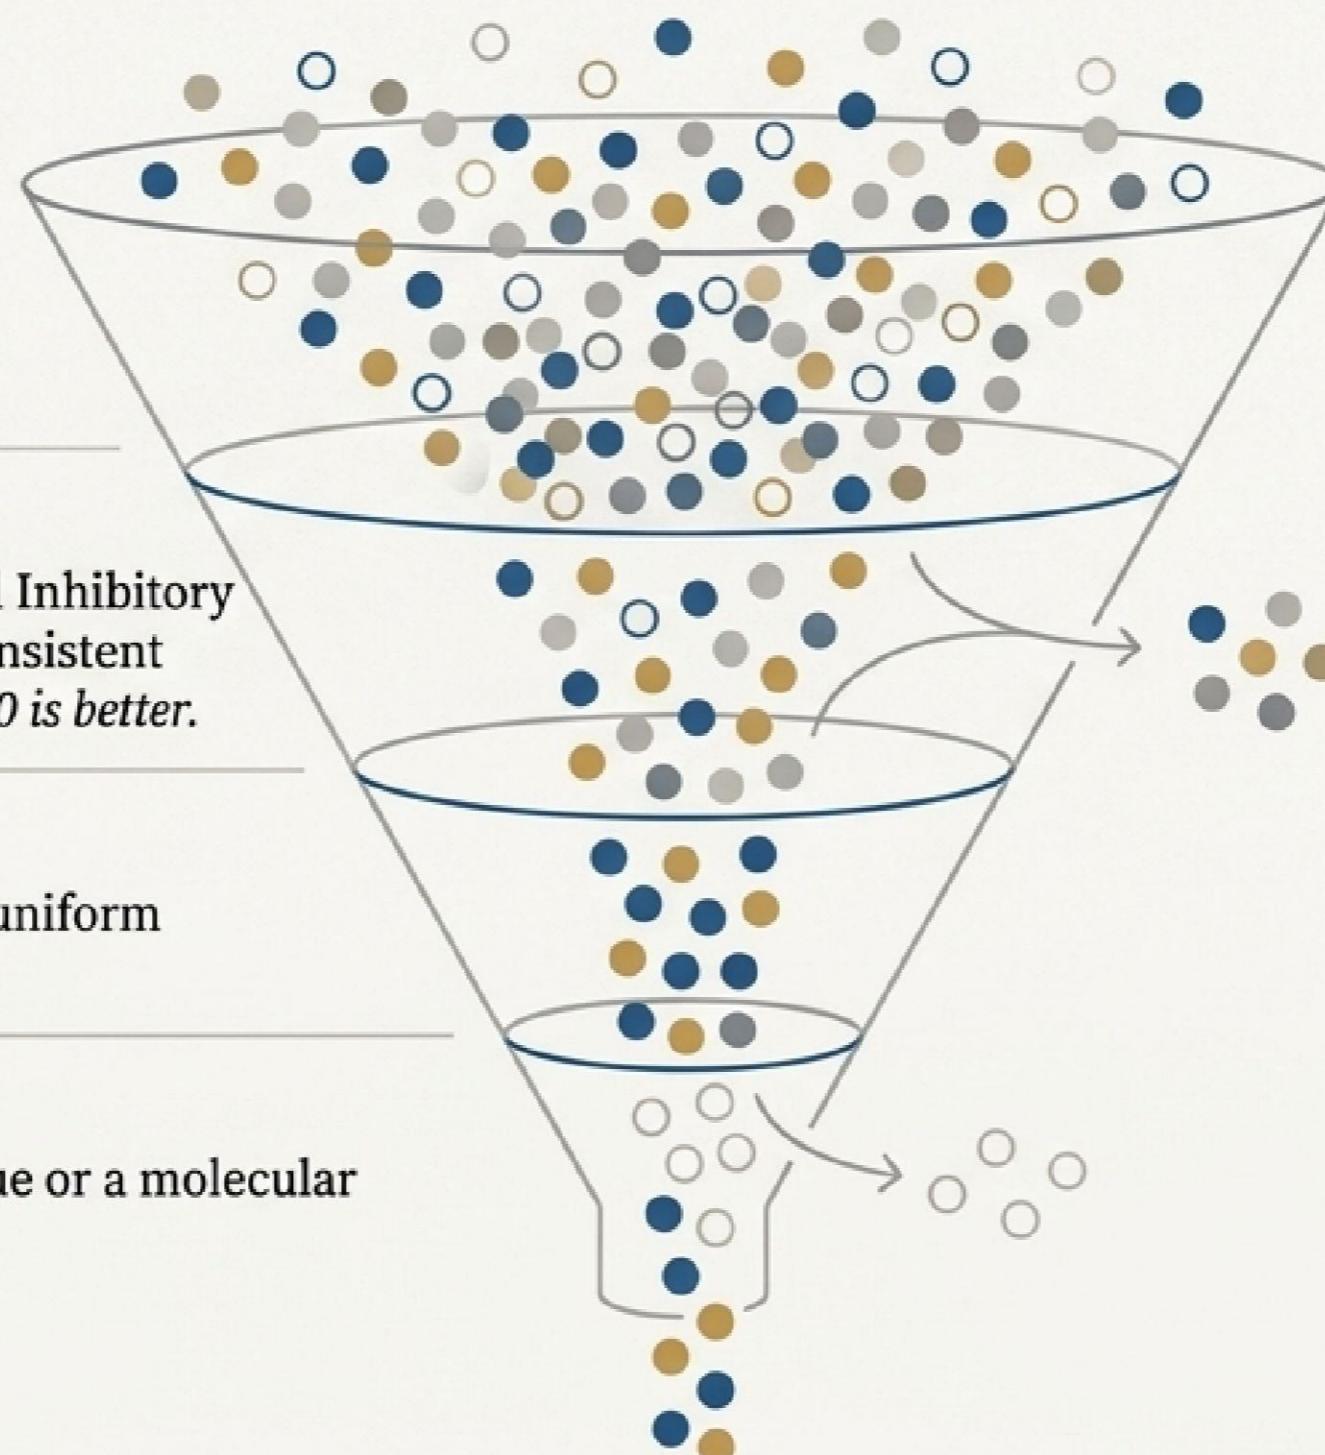
Keep only experiments using the 'Half Maximal Inhibitory Concentration' (IC50) metric. This ensures a consistent measure of potency of potency. *Note: Lower IC50 is better.*

Standardize Units

Convert all IC50 values to Nanomolar (nM) for uniform comparison.

Handle Missing Data

Remove any rows that lack a recorded IC50 value or a molecular structure (`canonical_smiles`).



Result: A high-quality, curated dataset ready for analysis.

Translating Chemistry into a Digital Language

1. Representing Structure: SMILES

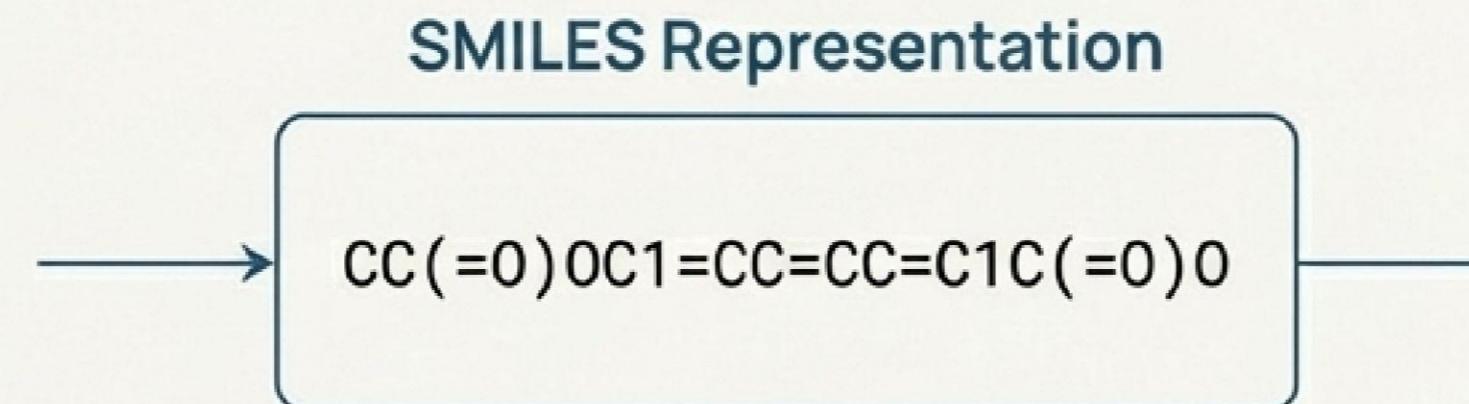
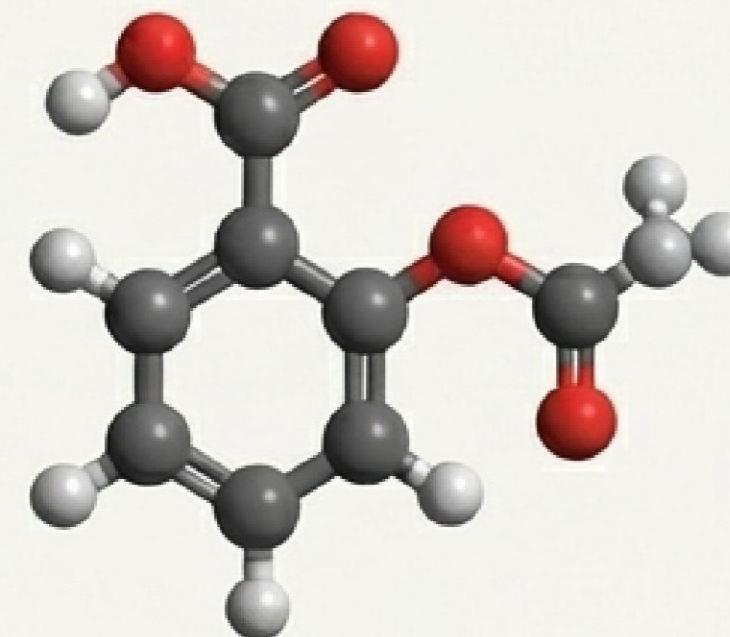
SMILES (Simplified Molecular Input Line Entry System) is a text-based notation that represents a complex 3D molecule as a simple string of characters.

Example: Water (H_2O) becomes `O`.
Ethanol ($\text{C}_2\text{H}_5\text{OH}$) becomes `CCO`.

2. Quantifying Properties: Lipinski's Rule of 5

A rule of thumb to evaluate ‘drug-likeness’ for oral medications. We calculate key descriptors:

- **Molecular Weight (MW):** < 500 Daltons
- **LogP (Solubility):** < 5
- **H-Bond Donors:** < 5
- **H-Bond Acceptors:** < 10

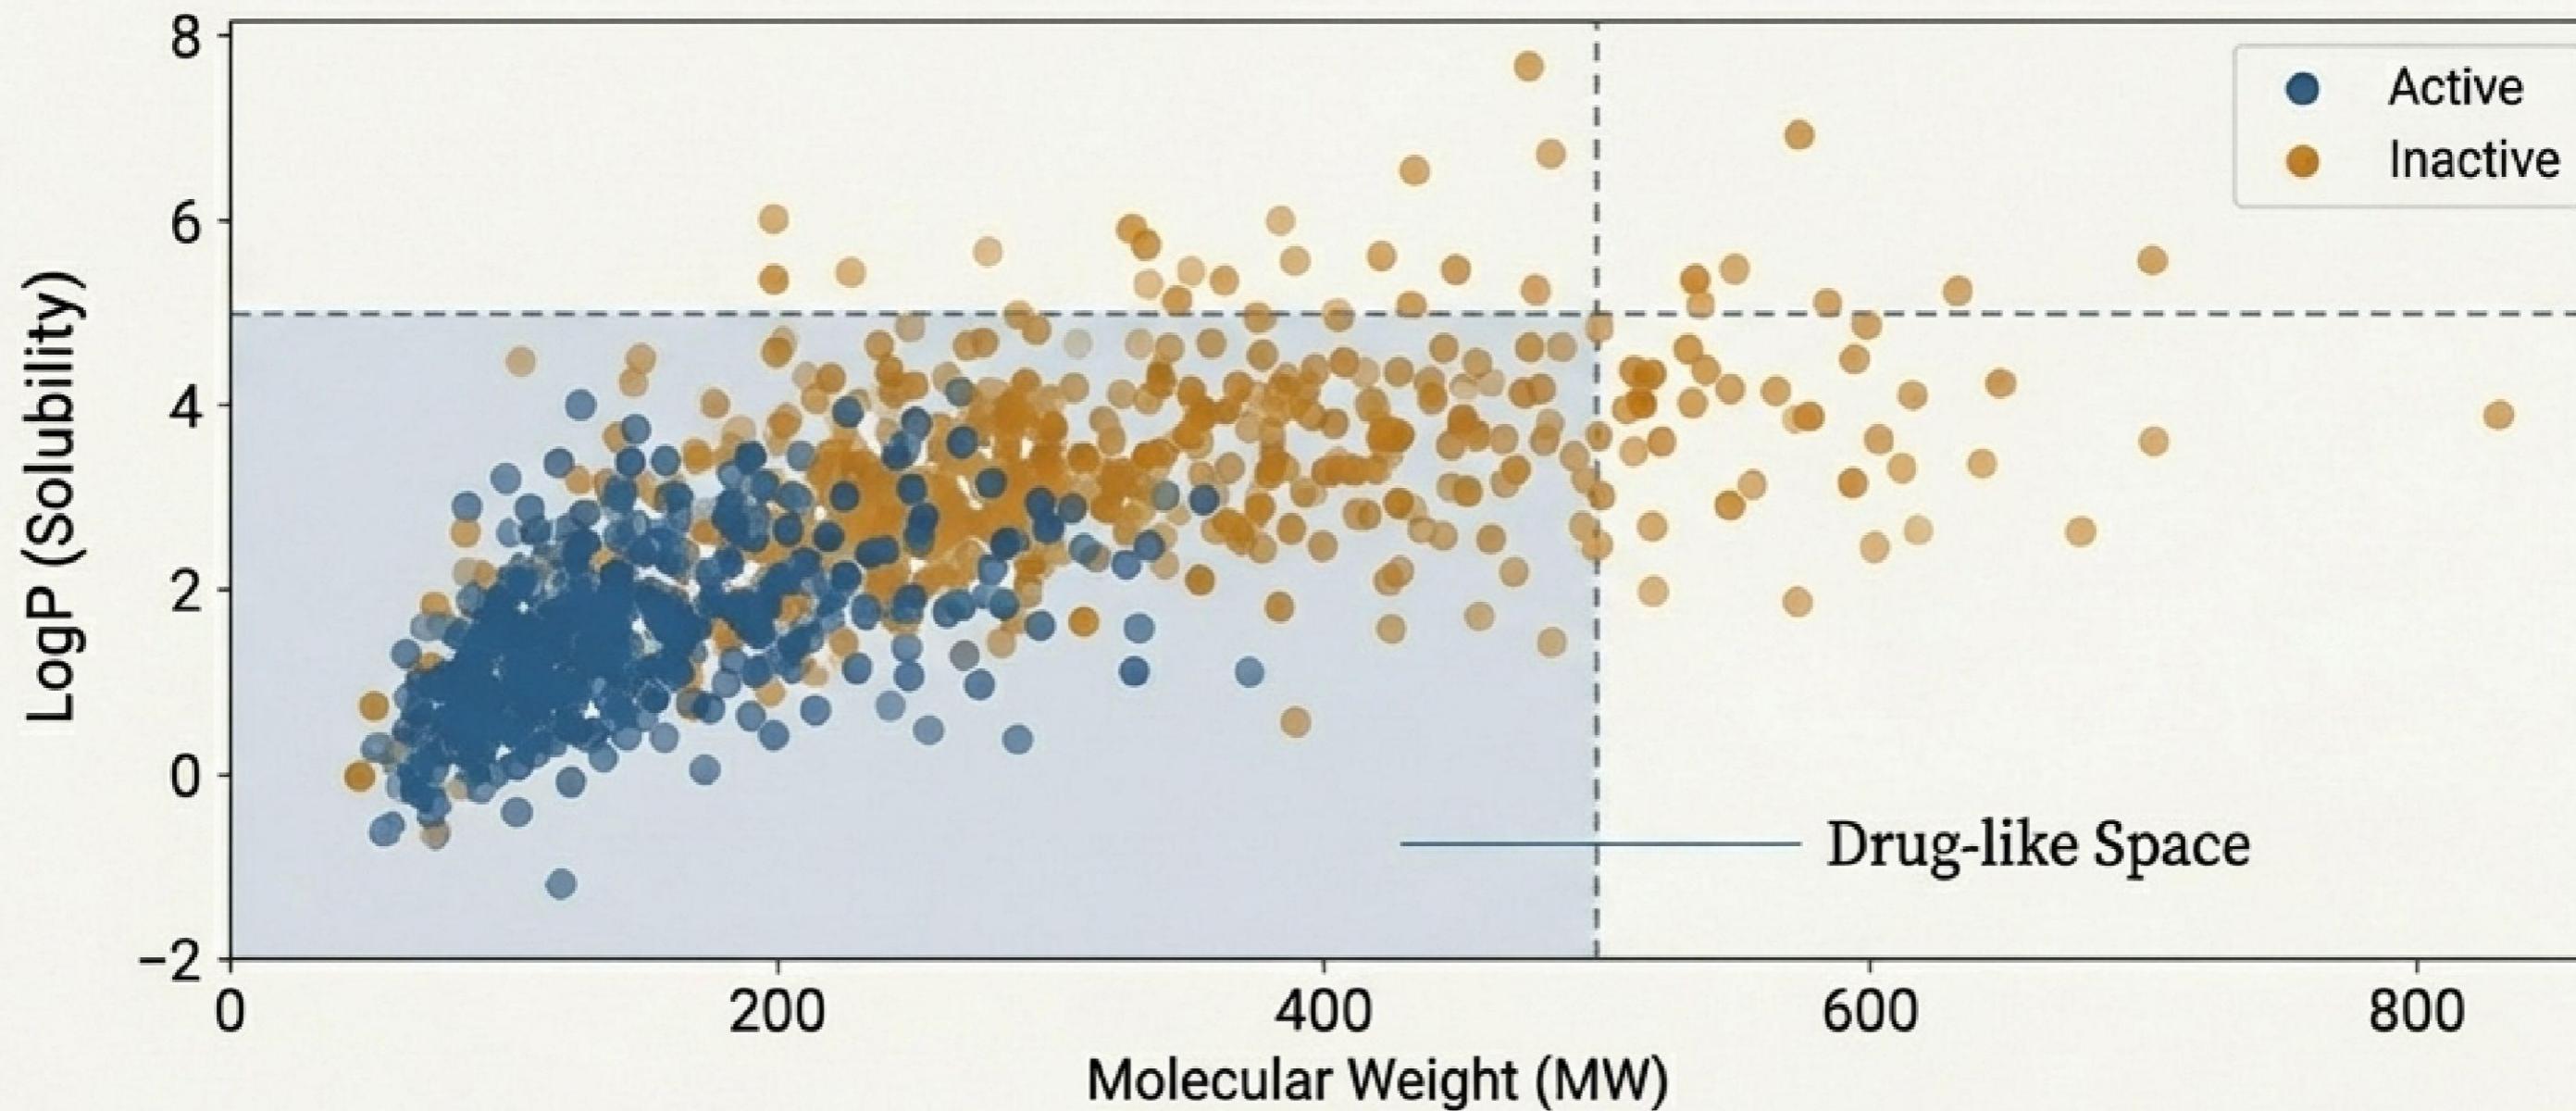


Lipinski Descriptors

Molecular Weight	180.16
LogP	1.19
H-Bond Donors	1
H-Bond Acceptors	4

Visualizing the Chemical Space of Active vs. Inactive Drugs

We plotted the Lipinski descriptors to see if potent ('Active') compounds occupy a different region of chemical space than weak ('Inactive') ones.



The plot shows a clear tendency for active compounds to cluster within the 'drug-like' boundaries defined by Lipinski's rules.

This visual difference was confirmed to be statistically significant using the Mann-Whitney U Test ($p < 0.05$).

The Challenge

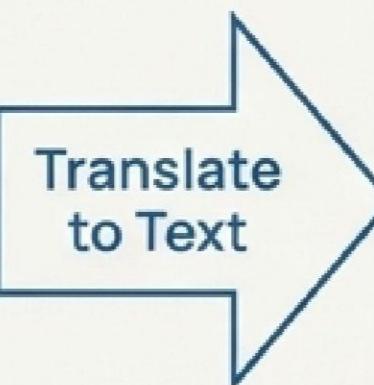
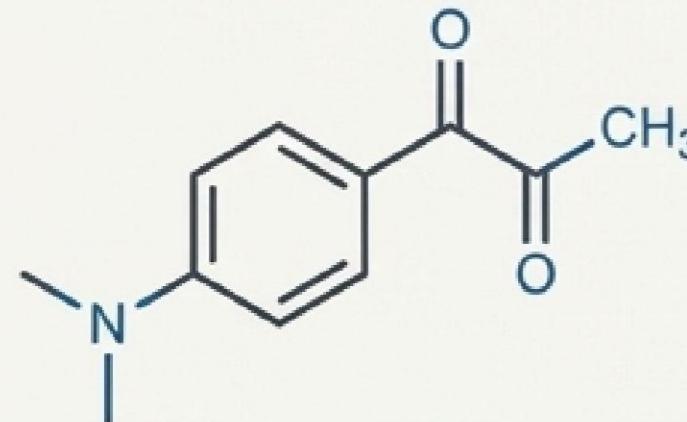
Machine learning algorithms can't process chemical diagrams or SMILES strings directly. They require numerical input.

Creating a 'Barcode' for Every Molecule

The Solution: Molecular Fingerprints

A molecular fingerprint is a binary vector (an array of 0s and 1s) that represents the presence or absence of specific chemical substructures within a molecule.

Chemical Structure

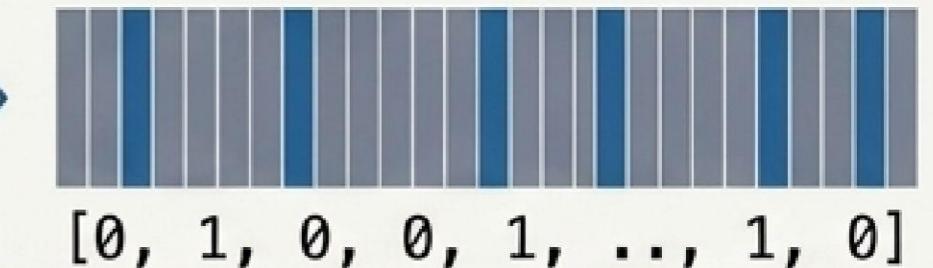


SMILES String

CC(=O)Nc1ccccc1C(=O)C(C)=O



881-bit PubChem Fingerprint



Our Method: PubChem Fingerprints

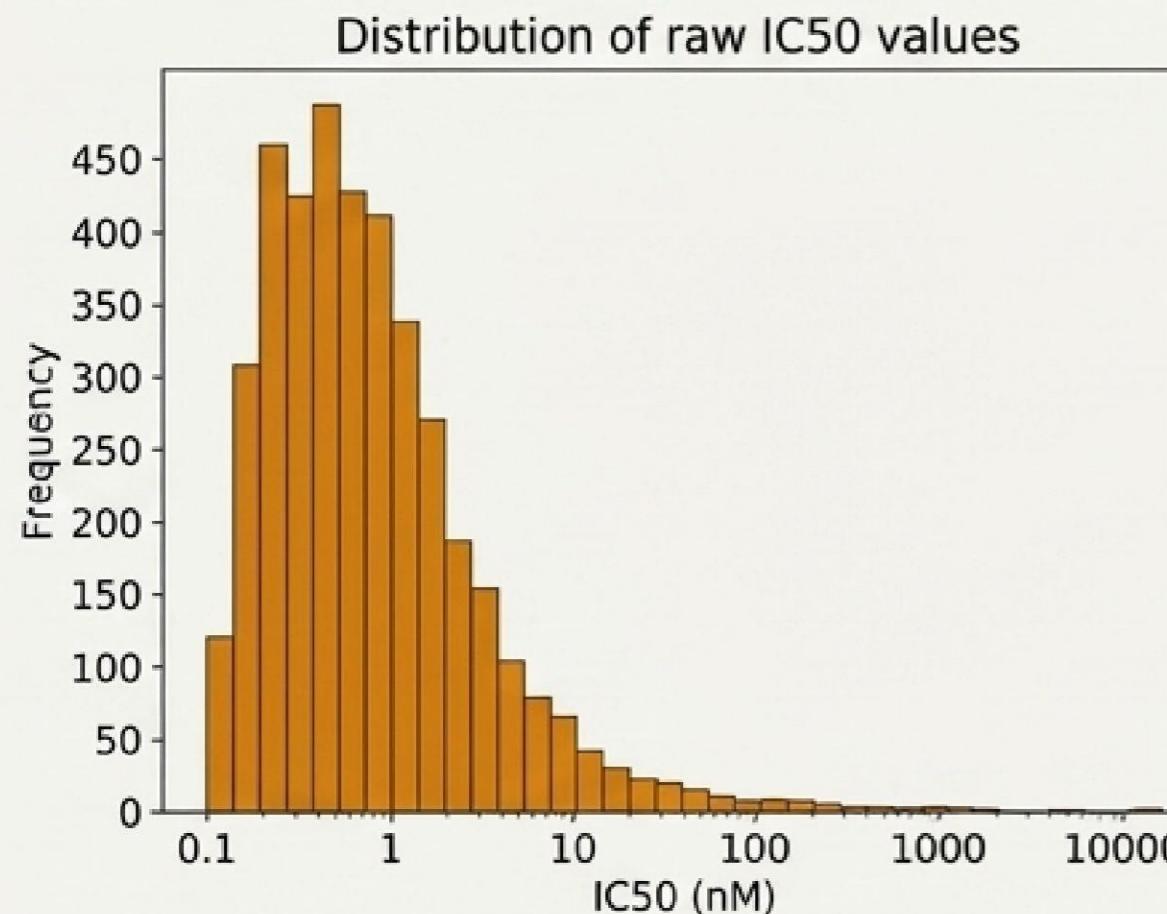
- We use a standard set of **881 chemical features**.
- For each molecule, we generate a vector of 881 bits. A '1' means the feature is present, and a '0' means it is absent.
- This vector becomes the unique, numerical identifier for the model to learn from.

Optimizing the Target Value: The pIC50 Transformation

Before training the model, we transform our target variable, IC50, into pIC50. This is essential for two reasons.

1. Normalizing the Distribution

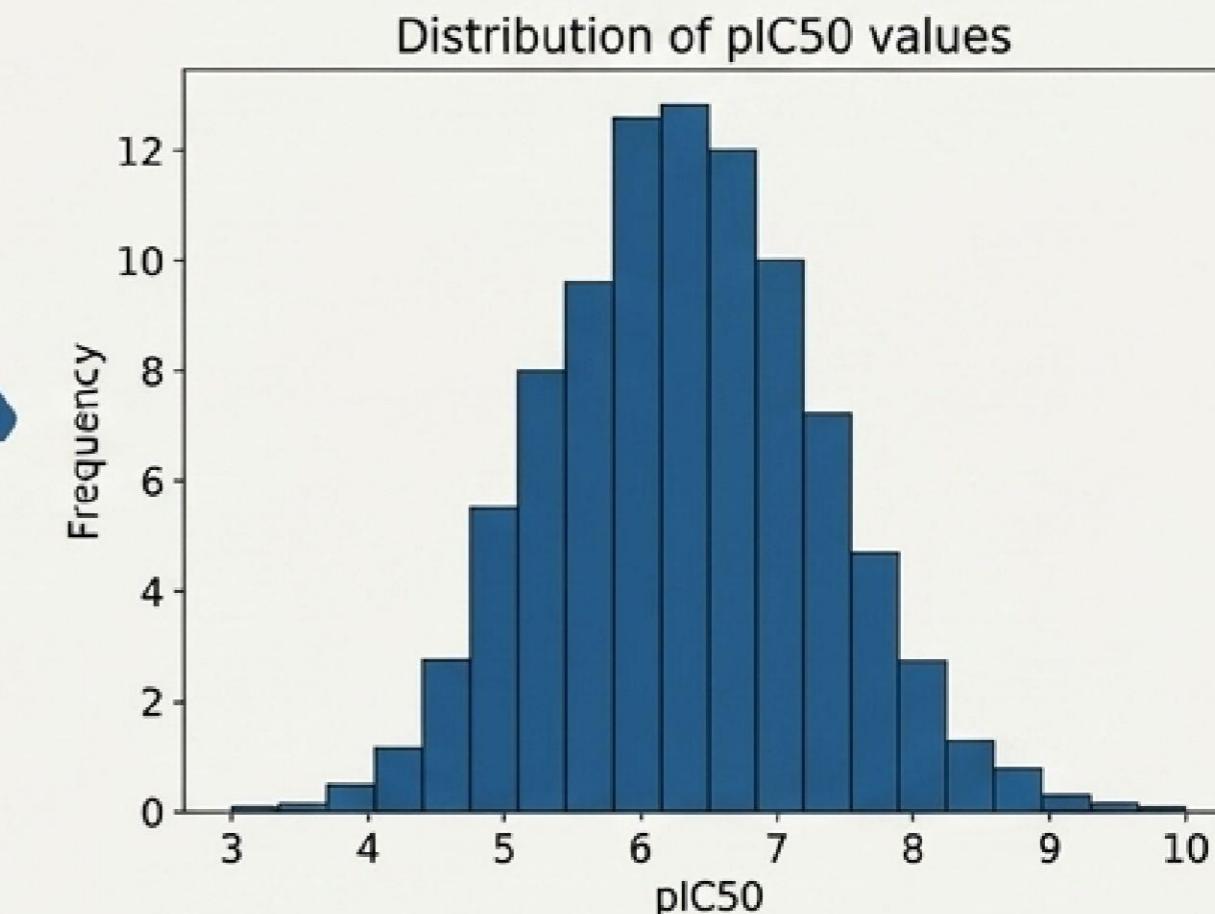
IC50 values are highly skewed, spanning many orders of magnitude. A log transform converts this skewed data into a more normal, linear distribution that is better suited for regression models.



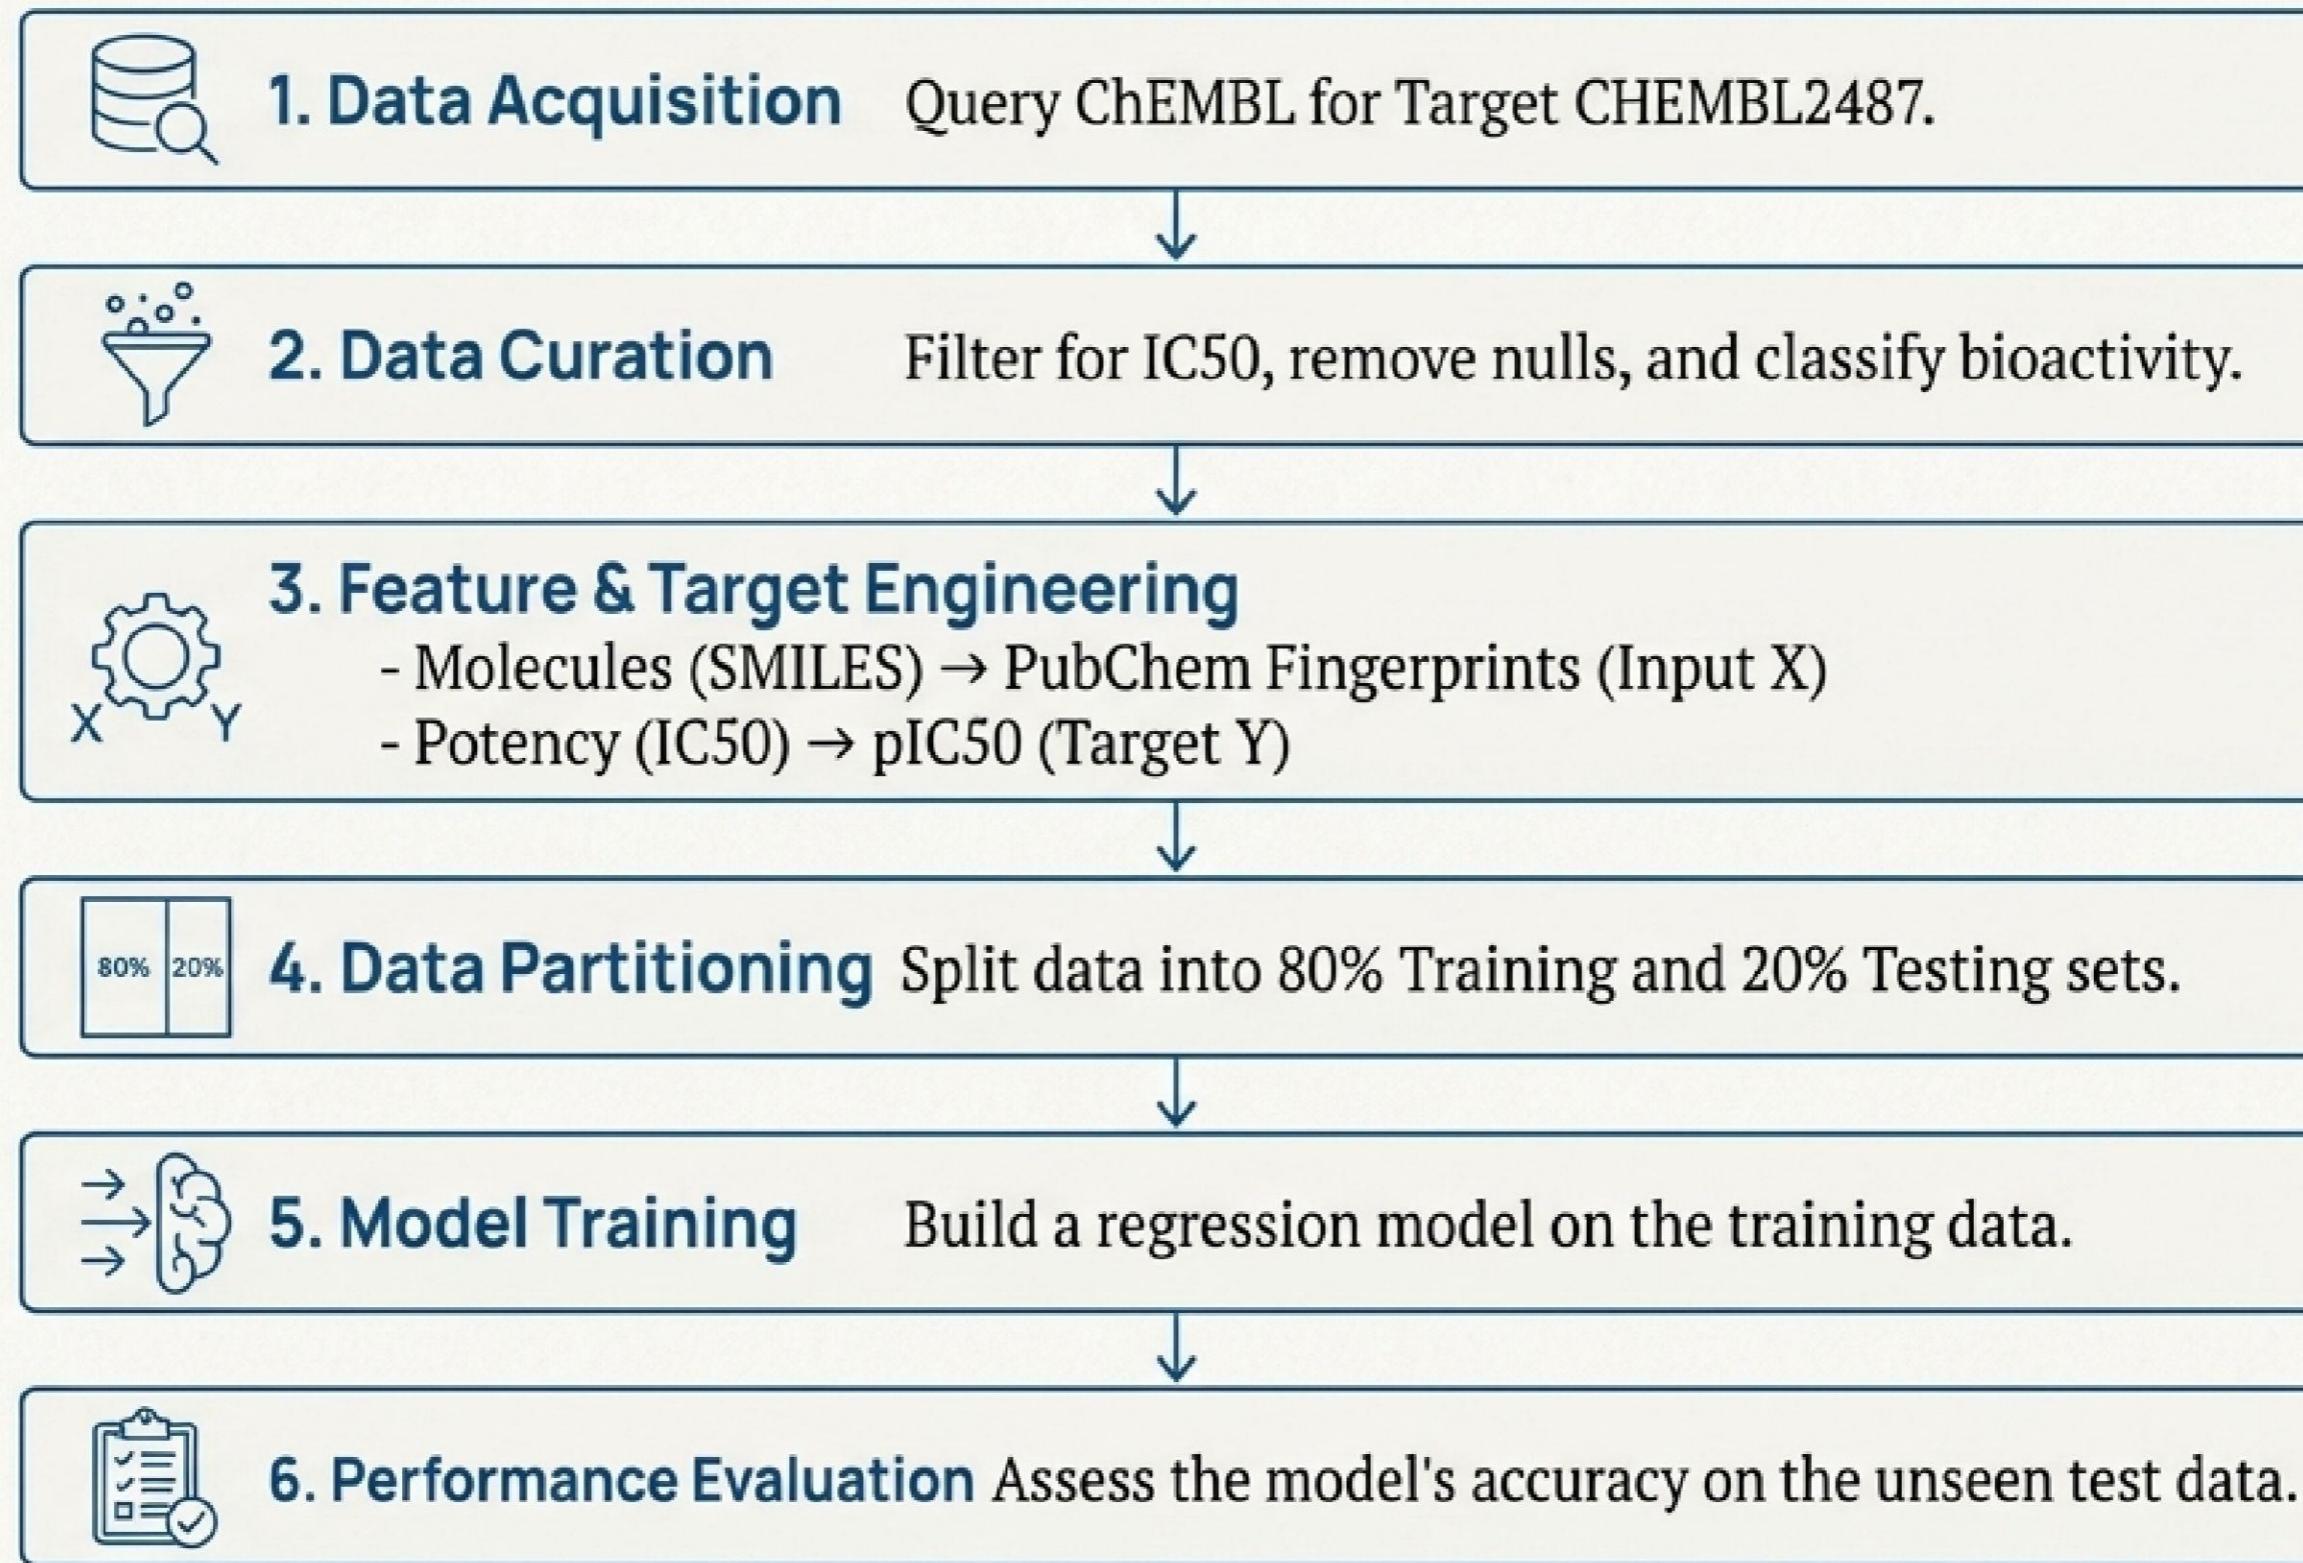
$$\text{pIC50} = -\log_{10}(\text{IC50})$$

2. Improving Interpretability

IC50 has an inverse scale (“Lower is Better”), which can be confusing. pIC50 ($-\log_{10}(\text{IC50})$) creates a direct scale where “**Higher is Better**,” which is more intuitive.



The End-to-End Computational Pipeline



Building and Training the Predictive Model

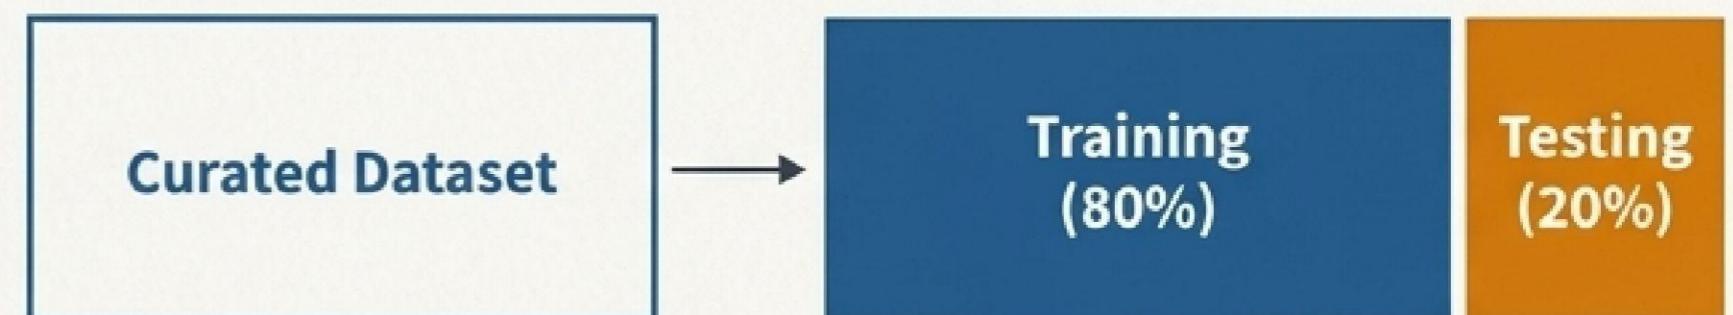
Step 1: Feature Selection

We first apply a `VarianceThreshold` to remove any of the 881 fingerprint features that are constant (all 0s or all 1s) across most of the dataset, as they provide no predictive information.



Step 2: Preventing Overfitting

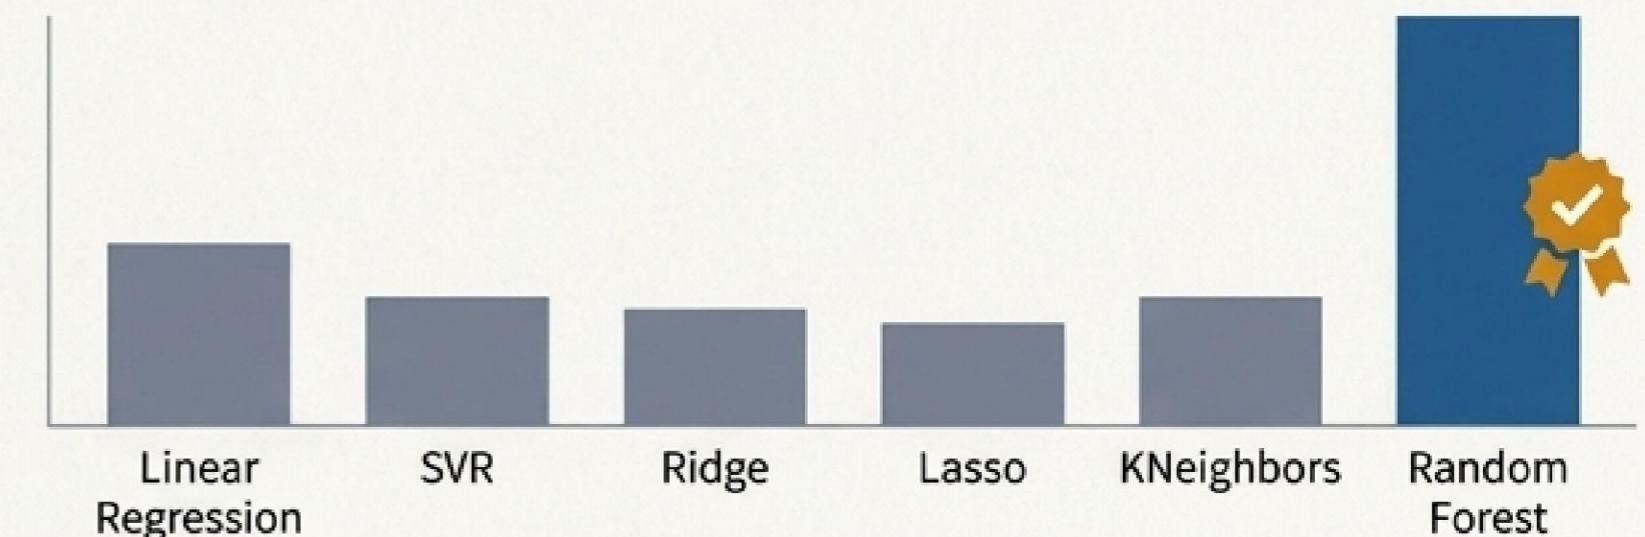
The curated dataset is split into an **80% Training Set** (for the model to learn from) and a **20% Testing Set** (held back for unbiased evaluation).



Step 3: Algorithm Selection

We used `LazyPredict`, a library that rapidly benchmarks dozens of regression algorithms. This “tournament” empirically identifies the best-performing model for our specific chemical data.

Winner: The **Random Forest Regressor** consistently emerged as a top performer.



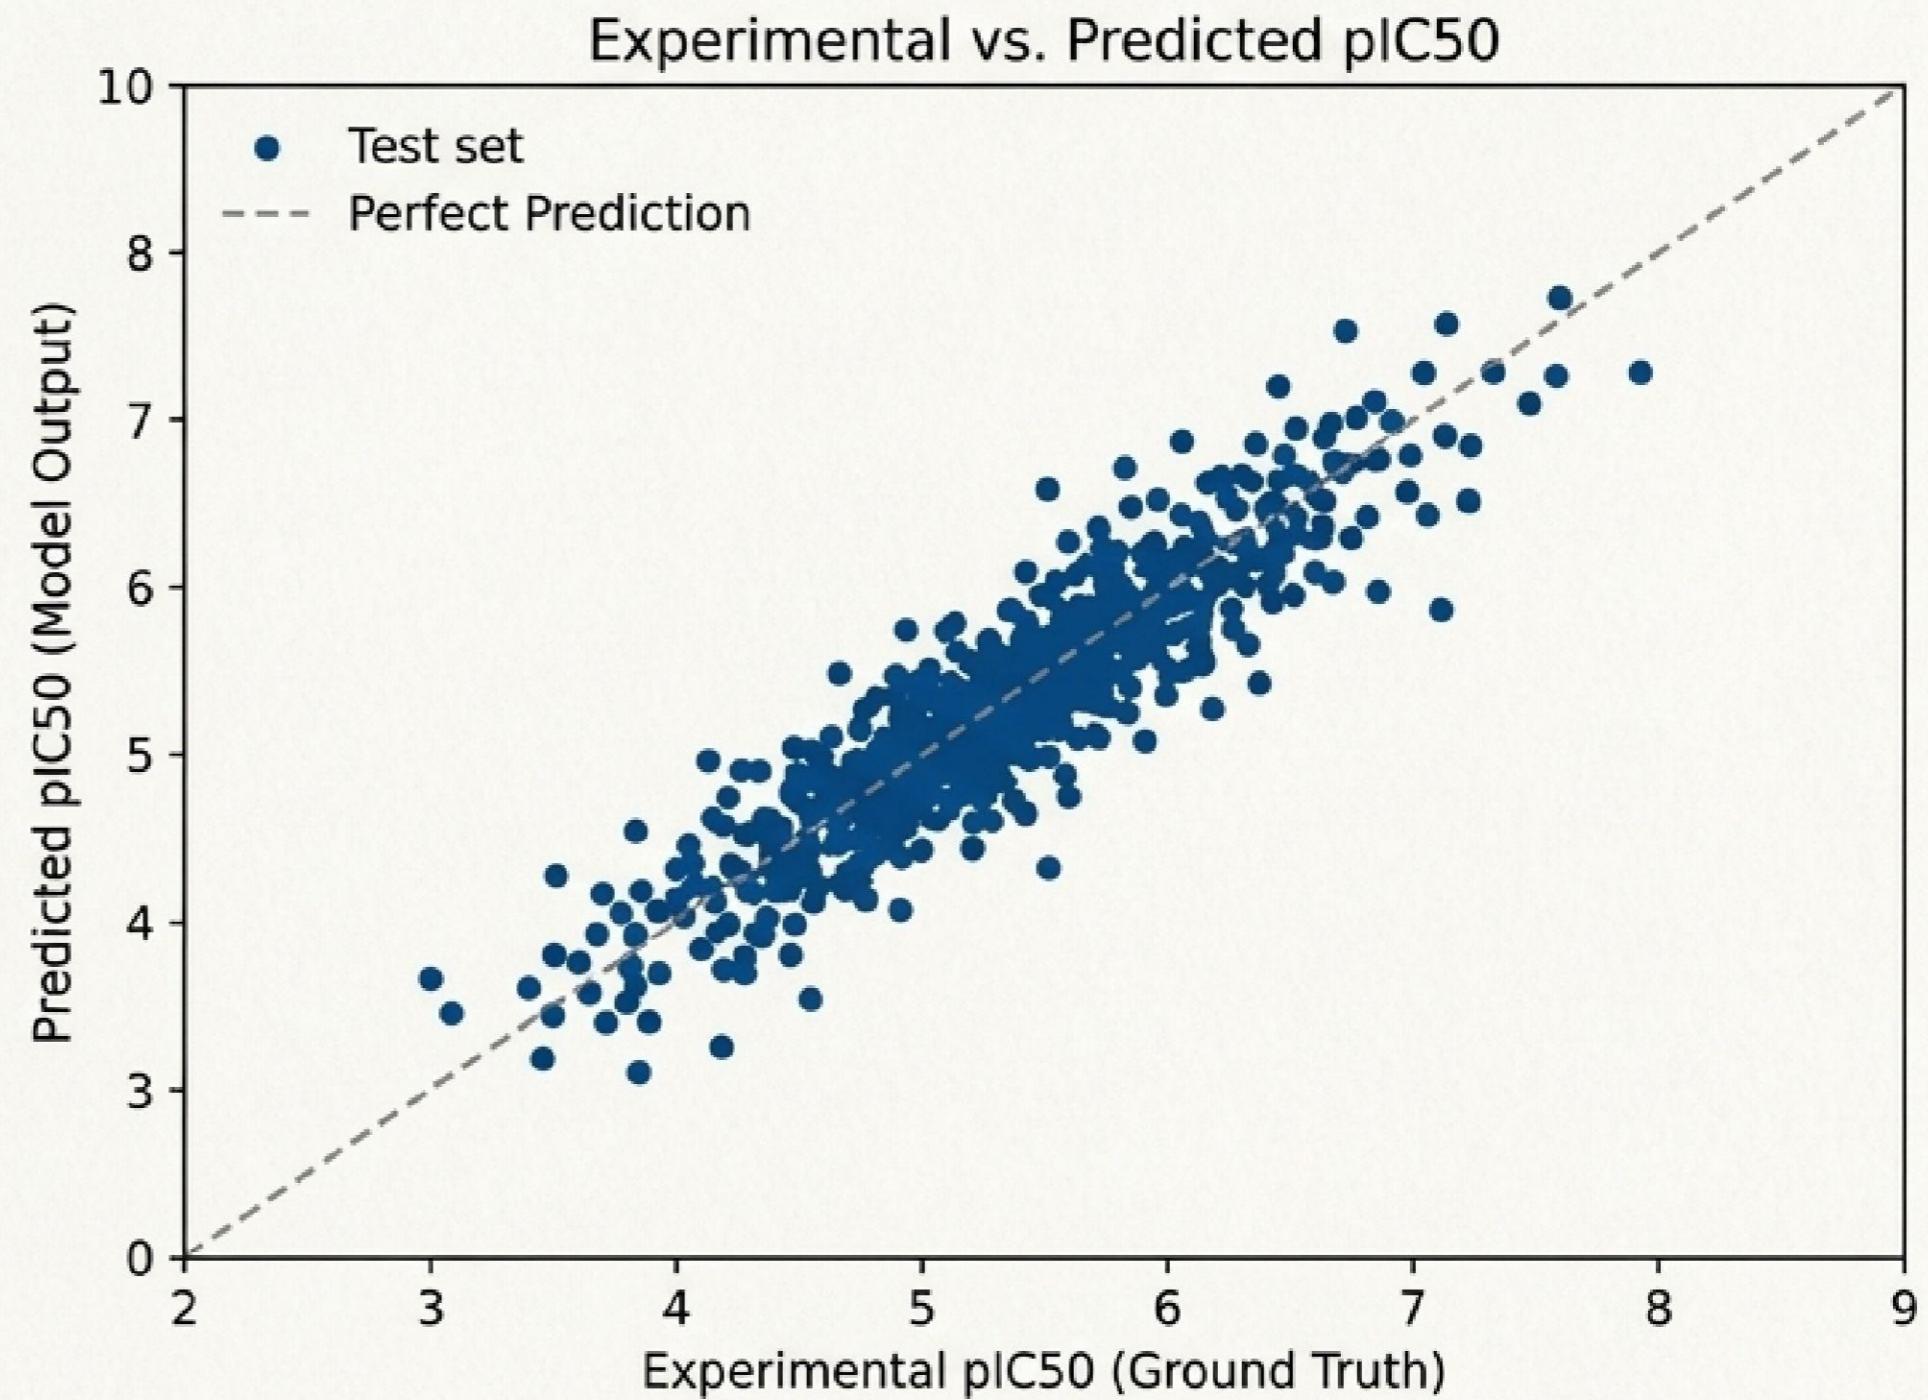
The Verdict: How Accurately Can We Predict Potency?

R² (Coefficient of Determination)

0.71

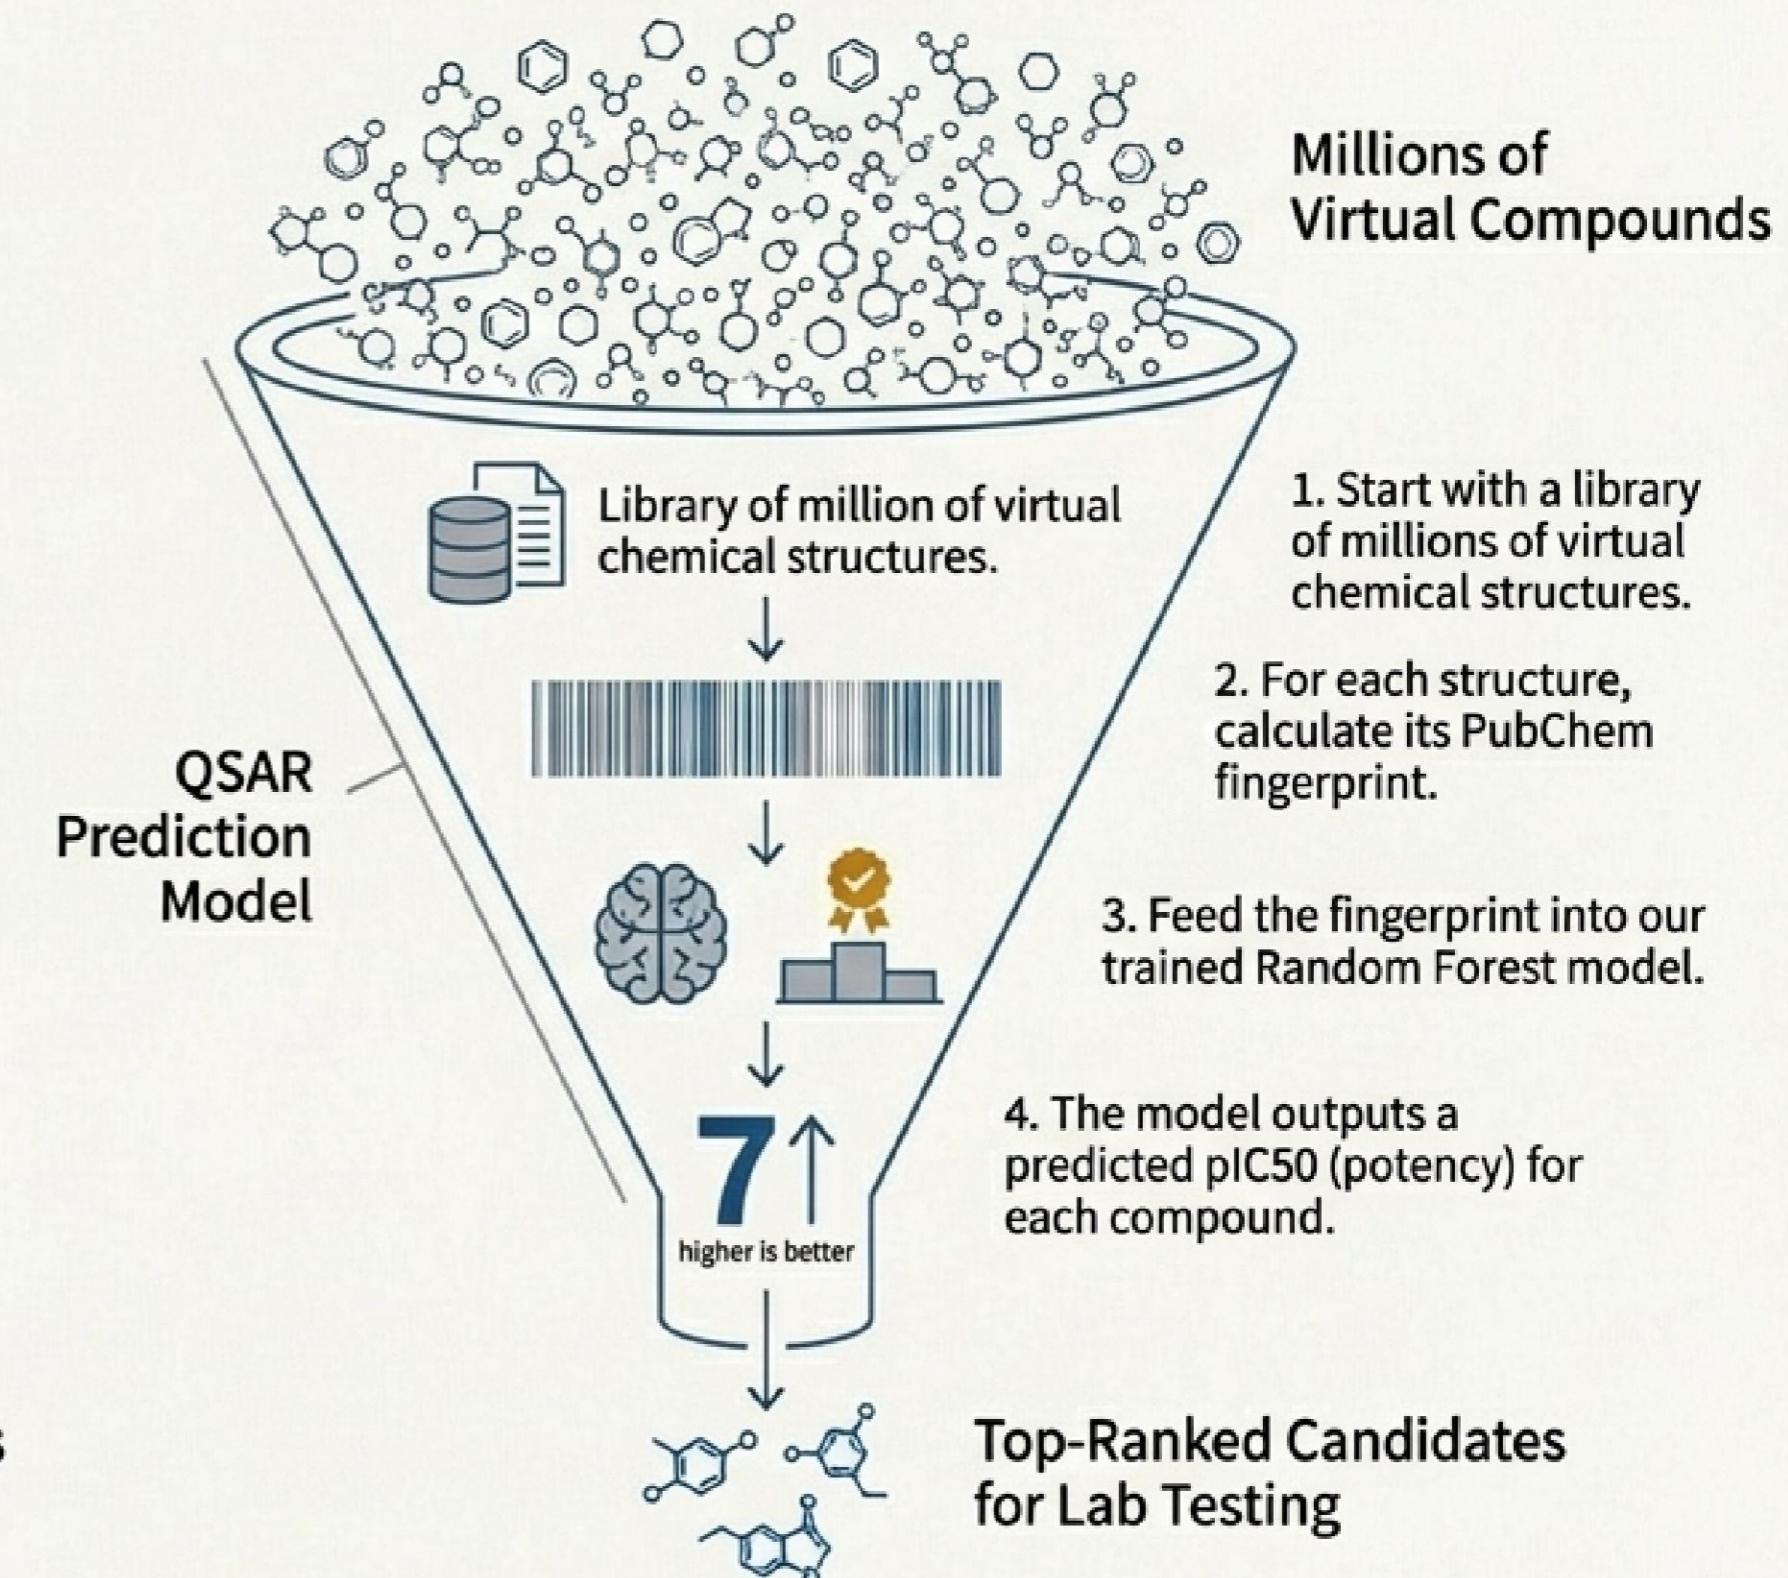
Interpretation

The R² value measures how much of the variance in the experimental data is explained by the model. A value of 0.71 is considered a strong result for QSAR models (where >0.6 is generally the goal), indicating that our model has successfully learned the underlying structure-activity relationships.



The Model in Action: Screening a Virtual Universe of Drugs

The true power of this model is not just predicting known compounds, but in its ability to screen massive virtual libraries of new, un-synthesized molecules.



Researchers can instantly identify and rank the most promising candidates, focusing expensive lab resources only on molecules with a high probability of success.

Accelerating the Quest for an Alzheimer's Cure

We began with a vast biological challenge and raw experimental data. By applying a rigorous data science pipeline, we successfully built and validated a QSAR model capable of predicting a compound's inhibitory potency against a key Alzheimer's target.

The Impact

This project is a case study in how computational methods can fundamentally alter the economics and timeline of drug discovery. By intelligently guiding lab research, machine learning can help accelerate our search for effective treatments for neurodegenerative diseases.

